

Satish Kallepalli Engineering Manager Resume

SATISH KUMAR KALLEPALLI

Engineering Manager | Platform & AI Products | Distributed Systems

San Francisco Bay Area | (510) 371-3621 | satishkallepalli940@gmail.com | linkedin.com/in/satish-k-28b227a7

PROFESSIONAL SUMMARY

Engineering Manager with 14+ years in enterprise software and hands-on people leadership building, scaling, and operating backend/platform teams for Fortune 500 organizations (Interpublic Group / Kinesso, Gap Inc.). Own hiring, performance management, career growth, delivery coaching, and cross-geo collaboration, with signature wins in team build, vendor-to-in-house migration, and production AI (Snowflake Cortex NLQ).

Proven at building teams from scratch, taking over vendor-delivered products, and migrating them onto in-house platforms and infrastructure without losing uptime or business continuity. Deliver inside IPG's Investment Console / reporting product area alongside Product, Data Engineering, and partner teams. Partner with Product, Data, Security, and Support on roadmaps, prioritization, and production ownership. Combine a strong technical bar (Java/Spring, cloud, data platforms, AI/LLM products) with a clear operating cadence: 1:1s, hiring loops, OKR and roadmap execution, incident leadership, and stakeholder communication.

Seeking Engineering Manager roles owning platform, data, or AI-enabled product engineering, where I can grow people, raise the delivery bar, and ship reliable systems at scale.

SELECTED LEADERSHIP IMPACT

- **People leadership:** Own hiring, onboarding, 1:1 coaching, performance reviews, growth plans, and succession for critical roles across product/platform engineering teams.
- **Product footprint:** Delivery inside Investment Console / reporting; partner across Product, Data Engineering, Security, and Support.
- **Team from scratch:** Recruited and built the Halogen engineering team end-to-end (roles, bar, onboarding, delivery rituals).
- **Vendor to in-house:** Led takeover of a vendor-operated media operations product and migration to in-house ownership and infrastructure, directing engineers across Mexico and Colombia while sustaining 99.9% uptime.
- **AI / NLQ:** Shipped entitlement-safe Snowflake Cortex Analyst chat in investment reporting: HARD entitlement injection, fail-closed SQL guardrails, rate limits, audit logging; cut ad-hoc planner-to-DE SQL requests.

- **Delivery outcomes:** Drove 80%+ reduction in manual rate-card processing; delivered multi-tenant investment reporting for IPG planners.
 - **Engineering excellence:** Cut deployment lead time ~40% (Jenkins to GitHub Actions; blue/green EKS); improved operability with Datadog and incident ownership (~35% user-satisfaction lift on Halogen).
 - **Scale:** Prior loyalty platforms at Gap processing millions of transactions/day.
-

CORE ENGINEERING MANAGER COMPETENCIES

Area	What I own
People & org	Hiring plans and loops, leveling, performance reviews, mentoring, career ladders, team health, cross-geo leadership
Delivery	Roadmap with Product, capacity planning, sprint/release cadence, risk and dependency management, stakeholder updates
Technical strategy	Architecture reviews, platform direction, quality bar, tech debt tradeoffs, AI/LLM productization with safety and auditability
Operations	SLOs/uptime, incident command, on-call health, observability, postmortems, production readiness
Vendor & transformation	Vendor transition, knowledge transfer, in-house migration, continuity plans, partner accountability
Cross-functional	Product / Data / Security / Support partnership; compliance-minded delivery (entitlements, audit logs)

TECHNICAL FOUNDATION

Languages: Java (14+ yrs), Python, TypeScript/JavaScript, SQL

Backend & platforms: Spring Boot, Kafka, AWS SQS/SNS, Redis, REST APIs, Node.js

Data: Snowflake (Semantic Views, Cortex Analyst), PostgreSQL, MongoDB, Oracle

Cloud & delivery: AWS (EKS, S3, SQS/SNS, Secrets Manager), Docker, Kubernetes, Terraform, GitHub Actions, Jenkins

AI products: Cortex Analyst, LangChain/LangGraph, RAG, prompt engineering, agentic workflows

Quality & ops: JUnit, Mockito, Gatling, ArchUnit, Datadog, Micrometer

PROFESSIONAL EXPERIENCE

Engineering Manager | Kinesso (Interpublic Group), San Francisco, CA

Feb 2022 – Present (*Engineering Manager scope: last 2 years; prior IC/Tech Lead contributions noted below*)

Scope: Manage product/platform engineering teams within IPG's Investment Console / reporting and media-ops product area. Own hiring, performance management, coaching, roadmap delivery, production reliability, and partnership with Product, Data Engineering, Security, and Support. Operate as a player-coach on architecture and critical delivery when needed.

People, process, and operating system

- Run the team operating cadence: 1:1s, staffing plans, interview loops, onboarding, yearly performance reviews, growth plans, and delivery coaching for tech leads and ICs.
- Set and defend engineering standards (design reviews, testing expectations, release readiness) so multiple workstreams ship without quality regressions (1,299+ automated tests kept green on key platforms).
- Translate Product strategy into executable roadmaps; manage capacity, dependencies, and risk; communicate status and tradeoffs to stakeholders.
- Build inclusive, high-accountability culture across US and LATAM (Mexico, Colombia) distributed teams.

Halogen Media Operations: team build + vendor to in-house migration | Feb 2024 – Present

- Built the Halogen engineering team from scratch (hiring plan, role definitions, bar-raising interviews, onboarding) to own IPG's flagship media operations platform (strategy through activation and billing).
- Took over a vendor-delivered product and led migration to in-house engineering ownership and infrastructure, including knowledge transfer, delivery continuity plans, and clear RACI with remaining vendor touchpoints.
- Directed distributed engineers through the transition while protecting business continuity and sustaining 99.9% uptime under high concurrent load.
- Modernized CI/CD (Jenkins to GitHub Actions), cutting deployment lead time ~40%; instituted blue/green EKS zero-downtime releases.
- Owned production operability (Datadog dashboards, alerting, incident leadership), contributing to ~35% improvement in user satisfaction.

AI Rate Card Automation | Mar 2025 – Present

- Led cross-functional delivery of a cloud-native AI platform that converts unstructured partner rate cards (Excel/CSV/PDF) into system-of-record data.
- Achieved 80%+ reduction in manual processing via async/concurrent extraction pipelines; drove partner iteration and post-incident accuracy improvements.
- Accountable for operability: caching, validation/retry, Snowflake + S3 + API integration, Datadog-backed production debugging.

AI-Powered NL Chat Assistant (Snowflake Cortex Analyst) | Jul 2026 – Present

- Designed and shipped an entitlement-safe NLQ interface embedded in IPG's investment reporting platform. Media planners ask business questions in English and get deterministic, audit-logged Snowflake answers via Cortex Analyst.

- Eliminated ad-hoc SQL requests from planners to Data Engineering by architecting an NL-to-SQL pipeline: HARD prompt injection of client/market entitlements, fail-closed SQL guardrail on Cortex output, and one-shot retry when required filters are missing.
- Built per-user and per-report token-bucket rate limiting (Bucket4j) with Micrometer/Datadog counters tagged by outcome and scope; Flyway-migrated chat_query_audit table capturing every request, applied filters, scope, and outcome for Support and regulatory review.
- Integrated Cortex Analyst against native Snowflake SEMANTIC VIEW objects (semantic model in the catalog, not runtime YAML); shipped report-scoped answer mode with UI toggle and applied-filter chips; 40 golden regression questions per dataset.
- Enforced package boundaries via an ArchUnit test suite so the chat module can later extract to a standalone service; kept a green 1,299-test unit suite through every refactor.

Investment Reporting Platform (GATE / Report Builder) | Jan 2024 – Present

- Drove delivery of a self-service multi-tenant reporting platform replacing Excel workflows over multi-billion-dollar media commitment and forecast data.
- Unblocked Data Engineering to onboard datasets in days instead of weeks via a metadata-driven query engine (5 datasets / one API).
- Owned reliability for large exports (100K+ rows): async bounded-concurrency pipelines, scheduled exports, schema evolution discipline (45+ Flyway migrations).

Outcome Planner | Feb 2022 – Jan 2024 (Tech Lead / Staff delivery leadership)

- Led distributed Spring Boot microservices with event-driven processing (SQS/SNS) for strategic media planning.
- Ran Agile delivery with bi-weekly releases; improved resilience through service decoupling and operational practices later scaled as EM.

Lead Software Engineer | Gap Inc., San Francisco, CA

May 2018 – Jan 2022

- Owned architecture and delivery leadership for Gap's enterprise loyalty platform serving millions of members.
- Designed high-throughput Kafka event-driven services processing millions of transactions/day; validated peak readiness with Gatling.
- Led zero-downtime monolith-to-microservices migration for customer profile services with data-integrity controls.
- Raised quality bar: automation that lifted coverage from 60% to 95%.
- Mentored engineers and influenced multi-team technical decisions.

Senior Software Engineer | Gap Inc. (via Nisum Technologies)

Dec 2011 – Apr 2018

- Led onsite/offshore engineering teams across retail platforms covering planning, delivery, and quality ownership.

- Delivered Price & Event Management and LOKI promotions engines at retail scale; cut regression time 60% via test automation.
-

EDUCATION & CERTIFICATIONS

B.Tech, Computer Science | Jawaharlal Nehru Technological University (JNTU), India | 2010

Certifications: Databricks Academy Agentic AI Development (Apr 2026) | Full Stack Developer | Agile Awareness for Teams and Scrum Masters